




PROBABILITY FOR STATISTICS AND DATA SCIENCE

Introduction to Probability: Cheat Sheet

Probability Formula | Sample Space | Expected Values | Complements

Words of welcome



You are here because you want to comprehend the basics of probability before you can dive into the world of statistics and machine learning. Understanding the driving forces behind key statistical features is crucial to reaching your goal of mastering data science. This way you will be able to extract important insight when analysing data through supervised machine learning methods like regressions, but also fathom the outputs unsupervised or assisted ML give you.

Bayesian Inference is a key component heavily used in many fields of mathematics to succinctly express complicated statements. Through Bayesian Notation we can convey the relationships between elements, sets and events. Understanding these new concepts will aid you in interpreting the mathematical intuition behind sophisticated data analytics methods.

Distributions are the main way we use to classify sets of data. If a dataset complies with certain characteristics, we can usually attribute the likelihood of its values to a specific distribution. Since many of these distributions have elegant relationships between certain outcomes and their probabilities of occurring, knowing key features of our data is extremely convenient and useful.

What is probability?

Probability is **the likelihood of an event occurring**. This event can be pretty much anything – getting heads, rolling a 4 or even bench pressing 225lbs. We measure probability with numeric values between 0 and 1, because we like to *compare* the relative likelihood of events. Observe the general probability formula.

$$P(X) = \frac{\textit{Preferred outcomes}}{\textit{Sample Space}}$$

Probability Formula:

- The Probability of event X occurring equals the *number* of preferred outcomes over the *number* of outcomes in the sample space.
- Preferred outcomes are the outcomes we want to occur or the outcomes we are interested in. We also call refer to such outcomes as “Favorable”.
- Sample space refers to all possible outcomes that can occur. Its “size” indicates the amount of elements in it.

If two events are independent:

The probability of them occurring simultaneously equals the product of them occurring on their own.

$$P(A \heartsuit) = P(A) \cdot P(\heartsuit)$$

Expected Values

Trial – Observing an event occur and recording the outcome.

Experiment – A collection of one or multiple trials.

Experimental Probability – The probability we assign an event, based on an experiment we conduct.

Expected value – the specific outcome we expect to occur when we run an experiment.

Example: Trial

Flipping a coin and recording the outcome.

Example: Experiment

Flipping a coin 20 times and recording the 20 individual outcomes.

In this instance, the **experimental probability** for getting heads would equal the number of heads we record over the course of the 20 outcomes, over 20 (the total number of trials).

The **expected value** can be numerical, Boolean, categorical or other, depending on the type of the event we are interested in. For instance, the expected value of the trial would be the more likely of the two outcomes, whereas the expected value of the experiment will be the number of time we expect to get either heads or tails after the 20 trials.

Expected value for **categorical** variables.

$$E(X) = n \times p$$

Expected value for **numeric** variables.

$$E(X) = \sum_{i=1}^n x_i \times p_i$$

Probability Frequency Distribution

What is a probability frequency distribution?:

A collection of the probabilities for each possible outcome of an event.

Why do we need frequency distributions?:

We need the probability frequency distribution to try and predict future events when the expected value is unattainable.

What is a frequency?:

Frequency is the number of times a given value or outcome appears in the sample space.

What is a frequency distribution table?:

The frequency distribution **table** is a table matching each distinct outcome in the sample space to its associated frequency.


How do we obtain the probability frequency distribution from the frequency distribution table?:

By dividing every frequency by the size of the sample space. (Think about the "favoured over all" formula.)

Sum	Frequency	Probability
2	1	$1/36$
3	2	$1/18$
4	3	$1/12$
5	4	$1/9$
6	5	$5/36$
7	6	$1/6$
8	5	$5/36$
9	4	$1/9$
10	3	$1/12$
11	2	$1/18$
12	1	$1/36$

Complements

The complement of an event is **everything** an event is **not**. We denote the complement of an event with an apostrophe.



Characteristics of complements:

- Can never occur simultaneously.
- Add up to the sample space. ($A + A' = \text{Sample space}$)
- Their probabilities add up to 1. ($P(A) + P(A') = 1$)
- The complement of a complement is the original event. ($((A')') = A$)

Example:

- Assume event A represents drawing a spade, so $P(A) = 0.25$.
- Then, A' represents **not** drawing a spade, so drawing a club, a diamond or a heart. $P(A') = 1 - P(A)$, so $P(A') = 0.75$.

Permutations

Permutations represent the number of different possible ways we can **arrange** a number of elements.

$$P(n) = n \times (n - 1) \times (n - 2) \times \cdots \times 1$$

The diagram illustrates the components of the permutation formula $P(n) = n \times (n - 1) \times (n - 2) \times \cdots \times 1$. Four teal arrows point from descriptive labels below to specific terms in the formula:

- An arrow points from "Permutations" to $P(n)$.
- An arrow points from "Options for who we put first" to n .
- An arrow points from "Options for who we put second" to $(n - 1)$.
- An arrow points from "Options for who we put last" to 1 .

Characteristics of Permutations:

- Arranging **all** elements within the sample space.
- No repetition.
- $P(n) = n \times (n - 1) \times (n - 2) \times \cdots \times 1 = n!$ (Called "n factorial")

Example:

- If we need to arrange 5 people, we would have $P(5) = 120$ ways of doing so.

Factorials

Factorials express the **product** of all integers from 1 to n and we denote them with the “!” symbol.

$$n! = n \times (n - 1) \times (n - 2) \times \cdots \times 1$$

Key Values:

- $0! = 1$.
- If $n < 0$, $n!$ does not exist.

Rules for factorial multiplication. (For $n > 0$ and $n > k$)

- $(n + k)! = n! \times (n + 1) \times \cdots \times (n + k)$
- $(n - k)! = \frac{n!}{(n - k + 1) \times \cdots \times (n - k + k)} = \frac{n!}{(n - k + 1) \times \cdots \times n}$
- $\frac{n!}{k!} = \frac{k! \times (k + 1) \times \cdots \times n}{k!} = (k + 1) \times \cdots \times n$

Examples: $n = 7$, $k = 4$

- $(7 + 4)! = 11! = 7! \times 8 \times 9 \times 10 \times 11$
- $(7 - 4)! = 3! = \frac{7!}{4 \times 5 \times 6 \times 7}$
- $\frac{7!}{4!} = 5 \times 6 \times 7$

Variations

Variations represent the number of different possible ways we can **pick** and **arrange** a number of elements.

Variations **with** repetition → $\bar{V}(n, p) = n^p$

Number of different elements available ↑

Number of elements we are arranging ←

Variations without repetition → $V(n, p) = \frac{n!}{(n-p)!}$

Number of different elements available ↑

Number of elements we are arranging ←

Intuition behind the formula. (With Repetition)

- We have n -many options for the first element.
- We **still have** n -many options for the second element because repetition is allowed.
- We have n -many options for each of the p -many elements.
- $n \times n \times n \dots n = n^p$

Intuition behind the formula. (Without Repetition)

- We have n -many options for the first element.
- We **only have** $(n-1)$ -many options for the second element because we cannot repeat the value for we chose to start with.
- We have less options left for each additional element.
- $n \times (n-1) \times (n-2) \dots (n-p+1) = \frac{n!}{(n-p)!}$

Combinations

Combinations represent the number of different possible ways we can pick a number of elements.

$$C(n, p) = \frac{n!}{(n - p)! p!}$$

Combinations

Total number of elements in the sample space

Number of elements we need to select

Characteristics of Combinations:

- Takes into account double-counting. (Selecting Johny, Kate and Marie is the same as selecting Marie, Kate and Johny)
- All the different permutations of a single combination are different variations.
- $C = \frac{V}{P} = \frac{n!/(n-p)!}{p!} = \frac{n!}{(n-p)!p!}$
- Combinations are symmetric, so $C_p^n = C_{n-p}^n$, since selecting p elements is the same as omitting n-p elements.

Combinations with separate sample spaces

Combinations represent the number of different possible ways we can pick a number of elements.

$$C = n_1 \times n_2 \times \cdots \times n_p$$

Diagram illustrating the formula for combinations with separate sample spaces:

- Combinations (points to C)
- Size of the first sample space. (points to n_1)
- Size of the second sample space. (points to n_2)
- Size of the last sample space. (points to n_p)

Characteristics of Combinations with separate sample spaces:

- The option we choose for any element does not affect the number of options for the other elements.
- The order in which we pick the individual elements is arbitrary.
- We need to know the size of the sample space for each individual element. ($n_1, n_2 \dots n_p$)

Winning the Lottery

To win the lottery, you need to satisfy two distinct independent events:

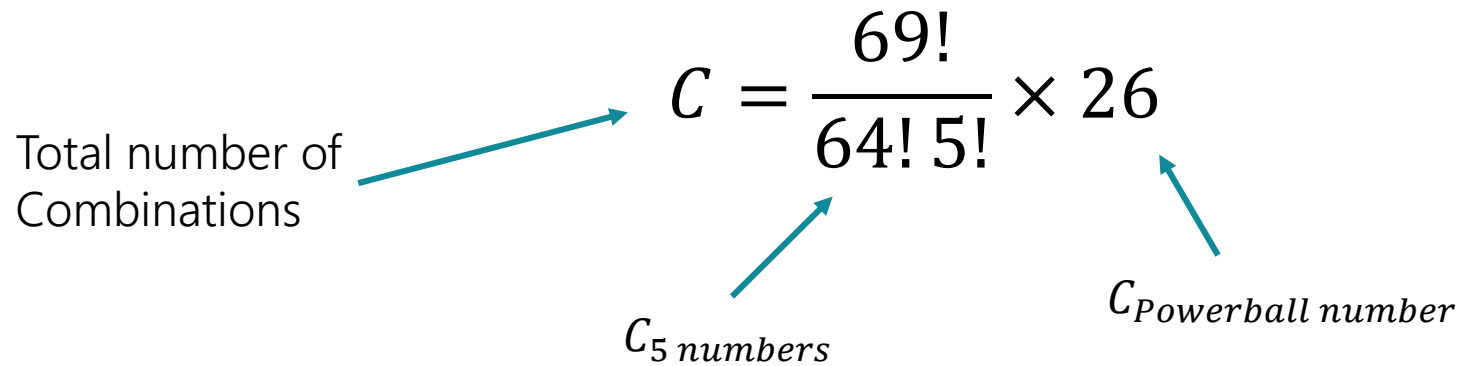
- Correctly guess the “Powerball” number. (From 1 to 26)
- Correctly guess the 5 regular numbers. (From 1 to 69)

Total number of Combinations

$$C = \frac{69!}{64! 5!} \times 26$$

$C_{5 \text{ numbers}}$

$C_{\text{Powerball number}}$

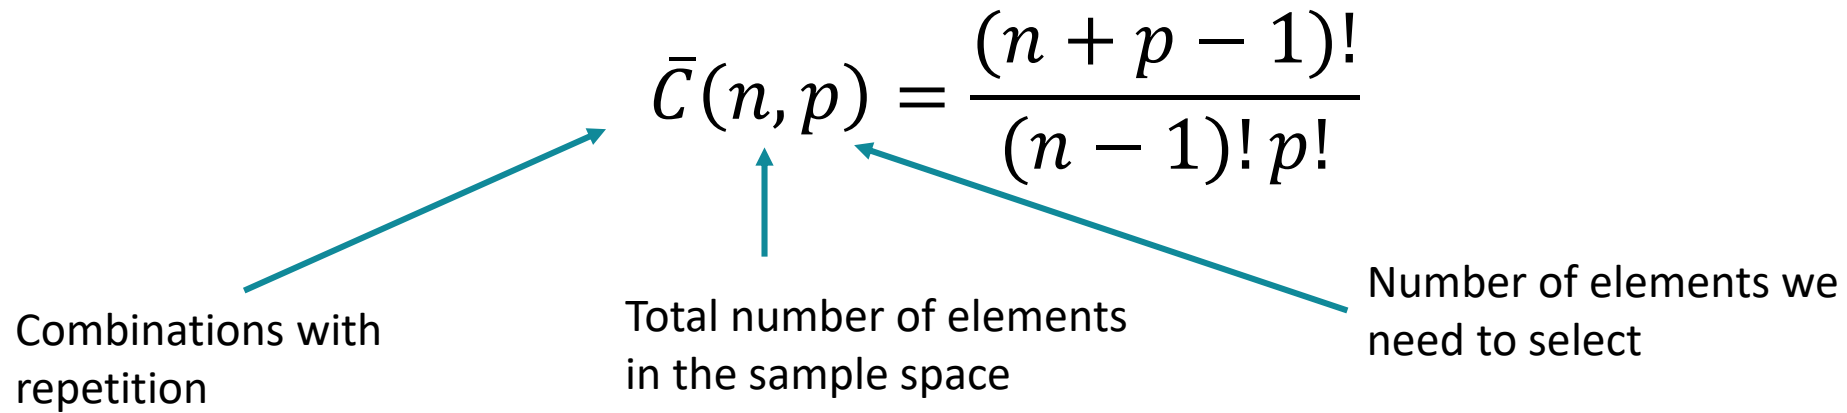


Intuition behind the formula:

- We consider the two distinct events as a combination of two elements with different sample sizes.
 - One event has a sample size of 26, the other has a sample size of C_5^{69} .
- Using the “favoured over all” formula, we find the probability of any single ticket winning equals $1 / (\frac{69!}{64! 5!} \times 26)$.

Combinations With Repetition

Combinations represent the number of different possible ways we can **pick** a number of elements. In special cases we can have repetition in combinations and for those we use a different formula.



Combinations with repetition

$$\bar{C}(n, p) = \frac{(n + p - 1)!}{(n - 1)! p!}$$

Total number of elements in the sample space

Number of elements we need to select

Now that you know what the formula looks like, we are going to walk you through the process of deriving this formula from the Combinations *without* repetition formula. This way you will be able to fully understand the intuition behind and not have to bother memorizing it.

Applications of Combinations with Repetition


To understand how combinations with repetition work you need to understand the instances where they occur.

We use combinations with repetition when the events we are dealing with, have sufficient quantity of each of the distinct values in the sample space.

One such example is the toppings on a pizza.

We can order extra of any given topping, so **repetition is allowed**. However, we **do not care about the order** in which the toppings are put on top of the pizza, so we cannot use variations.

Similar examples include picking the ice-cream flavours for a Sundae melt or the players for a Fantasy Football Team.



Pizza Example

To get a better grasp of the number of combinations we have, let us explore a specific example.

You are ordering a 3-topping pizza from your local pizza place but they only have 6 toppings left because it's late.

The toppings are as follows:
cheddar cheese, onions, green peppers, mushrooms, pepperoni and bacon.

Your pizza can have 3 different toppings, or you can repeat a topping up to 3 times.

You can either order a pizza with 3 different toppings, a pizza with 3 identical toppings or a pizza with 2 different toppings but having a double dose of one of them.



Methodology

The methodology we use for such combinations is rather abstract. We like to represent each type of pizza with a special sequence of 0s and 1s. To do so, we first need to select a specific order for the available ingredients.

We can reuse the order we wrote down earlier:
cheddar **c**heese, **o**nions, **g**reen peppers, **m**ushrooms, **p**epperoni and **b**acon.

For convenience we can refer to each ingredient by the associated letter we have highlighted (e.g "c" means cheese, and "o" means onions).

To construct the sequence for each unique type of pizza we follow 2 rules as we go through the ingredients in the order we wrote down earlier.

1. If we want no more from a certain topping, we write a **0** and **move to the next topping**.
2. If we want to include a certain topping, we write a **1** and **stay on the same topping**.
 - Not going to the next topping allows us to indicate if we want extra by adding another 1, before we move forward. Say, if we want our pizza to have extra cheese, the sequence would begin with "1, 1".
 - Also, we always apply rule 1 before moving on to another topping, so the sequence will actually start with "1, 1, 0".

Pizzas and Sequences

If we need to write a "0" after each topping, then every sequence would consist of 6 zeroes and 3 ones.

Let's look at some pizzas and the sequences they are expressed with.

A pizza with **cheese** and **extra peperoni** is represented by the sequence 1,0,0,0,0,1,1,0,0.

A vegan variety pizza with **onions**, **green peppers** and **mushrooms** would be represented by the sequence 0,1,0,1,0,1,0,0,0.

Now, what kind of pizza would the sequence 0,0,1,0,0,0,1,1,0 represent?

We can put the sequence into the table and see that it represents a cheese pizza with extra bacon.

C	O	G	M	P	B
1,0	0	0	0	1,1,0	0
0	1,0	1,0	1,0	0	0
0	0	1,0	0	0	1,1,0

Always Ending in 0

Notice how all the sequences we have examined end on a 0:

- 1,0,0,0,0,1,1,0,0
- 0,1,0,1,0,1,0,0,0
- 0,0,1,0,0,0,1,1,0

This is no coincidence, since according to rule 1 of our algorithm, we need to add a "0" at the end of the sequence, regardless of whether we wanted bacon or not.

That means that only the first 8 elements of the sequence can take different values.

Each pizza is characterized by the positions of the 3 "1"s in the sequence. Since only 8 of the positions in the sequence can take a value of "1", then the number of different pizzas would be the combination of any 3 of the 8 positions.

C	O	G	M	P	B
1,0	0	0	0	1,1,0	0
0	1,0	1,0	1,0	0	0
0	0	1,0	0	0	1,1,0

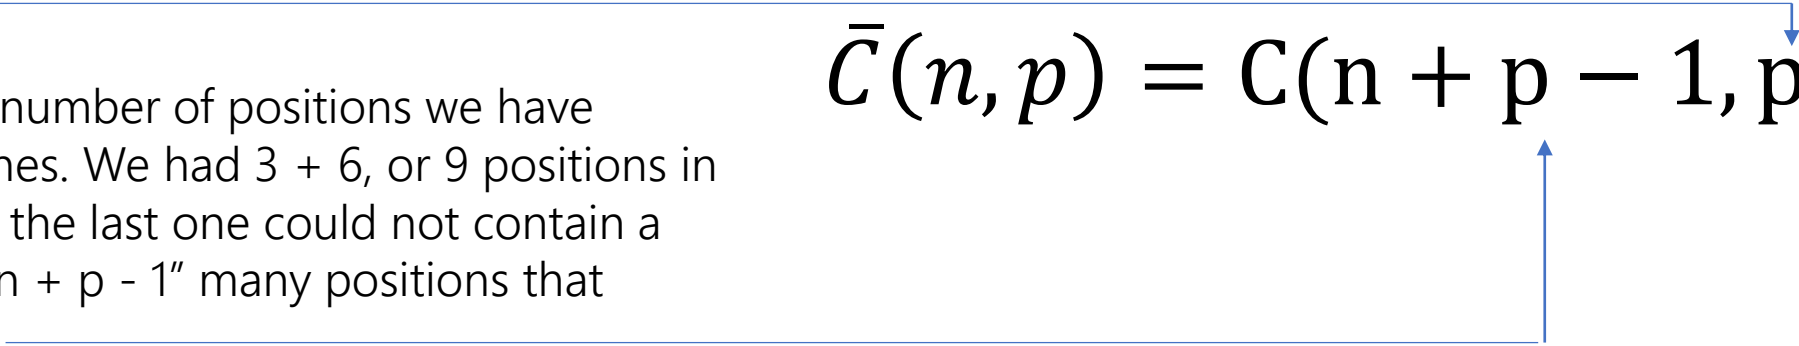
Positions

As stated before, we have 3 "1s" and 8 different positions. Therefore, the number of pizzas we can get would be the number of combinations of picking 3 elements out of a set of 8. This means we can transform combinations **with** repetition to combinations **without** repetition.

$$\bar{C}(6,3) = C(8,3)$$

Let's observe the values 3 and 8 for a moment and try to generalize the formula. "3" represents the amount of toppings we need to pick, so it is still equivalent to p.

"8" represents the number of positions we have available for the ones. We had 3 + 6, or 9 positions in total, but we knew the last one could not contain a "1". Thus, we had "n + p - 1" many positions that could contain a 1.


$$\bar{C}(n, p) = C(n + p - 1, p)$$

The Final Step

Now that we know the relationship between the number of combinations **with** and **without** repetition, we can plug in “ $n+p-1$ ” into the combinations without repetition formula to get:

$$\bar{C}(n, p) = C(n + p - 1, p) = \frac{(n + p - 1)!}{((n + p - 1) - p)! p!} = \frac{(n + p - 1)!}{(n - 1)! p!}$$

This is the exact same formula we showed you at the beginning.

Before we continue to the next lecture, let's make a quick recap of the algorithm and the formula.

1. We started by ordering the possible values and expressing every combination as a sequence.
2. We examined that only certain elements of the sequence may differ.
3. We concluded that every unique sequence can be expressed as a combination of the positions of the “1” values.
4. We discovered a relationship between the formulas for combinations with and without repetition.
5. We used said relationship to create a general formula for combinations with repetition.

Symmetry of Combinations

Let's see the algebraic proof of the notion that selecting p -many elements out of a set of n is the same as omitting $n-p$ many elements.

For starters, recall the combination formula:

$$C(n, p) = \frac{n!}{(n-p)! p!}$$

If we plug in $n-p$ for p , we get the following:


$$C(n, n-p) = \frac{n!}{(n-(n-p))! (n-p)!} = \frac{n!}{(n-n+p)! (n-p)!} = \frac{n!}{p! (n-p)!} = \frac{n!}{(n-p)! p!} = C(n, p)$$

Therefore, we can conclude that $C(n, p) = C(n, n-p)$.

Bayesian Notation

A **set** is a collection of elements, which hold certain values. Additionally, every event has a set of outcomes that satisfy it.

The **null-set** (or **empty set**), denoted \emptyset , is an set which contain no values.



Notation:

$x \in A$

$A \ni x$

$x \notin A$

$\forall x$:

$A \subseteq B$

Interpretation:

"Element x is a part of set A."

"Set A contains element x."

"Element x is NOT a part of set A."

"For all/any x such that..."

"A is a subset of B"

Example:

$2 \in \text{All even numbers}$

$\text{All even numbers} \ni 2$

$1 \notin \text{All even numbers}$

$\forall x: x \in \text{All even numbers}$

$\text{Even numbers} \subseteq \text{Integers}$

Remember! Every set has at least 2 subsets.

- $A \subseteq A$
- $\emptyset \subseteq A$

Multiple Events

The sets of outcomes that satisfy two events A and B can interact in one of the following 3 ways.


Not touch at all.



Intersect (Partially Overlap)



One completely overlaps the other.



Examples:

A -> Diamonds

B -> Hearts

Diamonds


Queens

Red Cards

Diamond

Intersection

The **intersection** of two or more events expresses the set of outcomes that satisfy all the events simultaneously. Graphically, this is the area where the sets intersect.




We denote the intersection of two sets with the “intersect” sign, which resembles an upside-down capital letter U:

$$A \cap B$$


Union

The **union** of two or more events expresses the set of outcomes that satisfy at least one of the events. Graphically, this is the area that includes both sets.



Mutually Exclusive Sets

Sets with no overlapping elements are called **mutually exclusive**. Graphically, their circles never touch.



If $A \cap B = \emptyset$, then the two sets are mutually exclusive.

Remember:

All complements are mutually exclusive, but not all mutually exclusive sets are complements.

Example:

Dogs and Cats are mutually exclusive sets, since no species is simultaneously a feline and a canine, but the two are not complements, since there exist other types of animals as well.

Independent and Dependent Events

If the likelihood of event A occurring ($P(A)$) is affected event B occurring, then we say that A and B are **dependent** events. Alternatively, if it isn't – the two events are **independent**.

We express the probability of event A occurring, given event B has occurred the following way **$P(A|B)$** . We call this the conditional probability.

Independent:

- All the probabilities we have examined so far.
- The outcome of A does not depend on the outcome of B.
- $P(A|B) = P(A)$

Example

- A -> Hearts
- B -> Jacks

Dependent

- New concept.
- The outcome of A depends on the outcome of B.
- $P(A|B) \neq P(A)$

Example

- A -> Hearts
- B -> Red

Conditional Probability

For any two events A and B, such that the likelihood of B occurring is greater than 0 ($P(B) > 0$), the conditional probability formula states the following.

Probability of A,
given B has
occurred

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

Probability of
the intersection.

Probability of
event B

Intuition behind the formula:

- Only interested in the outcomes where B is satisfied.
- Only the elements in the intersection would satisfy A as well.
- Parallel to the “favoured over all” formula:
 - Intersection = “preferred outcomes”
 - B = “sample space”

Remember:

- Unlike the union or the intersection, changing the order of A and B in the conditional probability alters its meaning.
- $P(A|B)$ is not the same as $P(B|A)$, even if $P(A|B) = P(B|A)$ numerically.
- The two conditional probabilities possess **different meanings** even if they have equal values.

Law of total probability

The **law of total probability** dictates that for any set A , which is a union of many mutually exclusive sets B_1, B_2, \dots, B_n , its probability equals the following sum.

$$P(A) = P(A|B_1) \times P(B_1) + P(A|B_2) \times P(B_2) + \dots + P(A|B_n) \times P(B_n)$$

Conditional Probability of A , given B_1 has occurred.

Probability of B_1 occurring.

Conditional Probability of A , given B_2 has occurred.

Probability of B_2 occurring.

Intuition behind the formula:

- Since $P(A)$ is the union of mutually exclusive sets, so it equals the **sum of the individual sets**.
- The **intersection** of a union and one of its subsets is the entire subset.
- We can rewrite the conditional probability formula $P(A|B) = \frac{P(A \cap B)}{P(B)}$ to get $P(A \cap B) = P(A|B) \times P(B)$.
- Another way to express the law of total probability is:
 - $P(A) = P(A \cap B_1) + P(A \cap B_2) + \dots + P(A \cap B_n)$


Additive Law

The additive law calculates the probability of the union based on the probability of the individual sets it accounts for.

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$



Probability of the
union



Probability of
the intersection

Intuition behind the formula

- Recall the formula for finding the size of the Union using the size of the Intersection:
 - $A \cup B = A + B - A \cap B$
- The probability of each one is simply its size over the size of the sample space.
- This holds true for any events A and B.

The Multiplication Rule

The multiplication rule calculates the probability of the intersection based on the conditional probability.

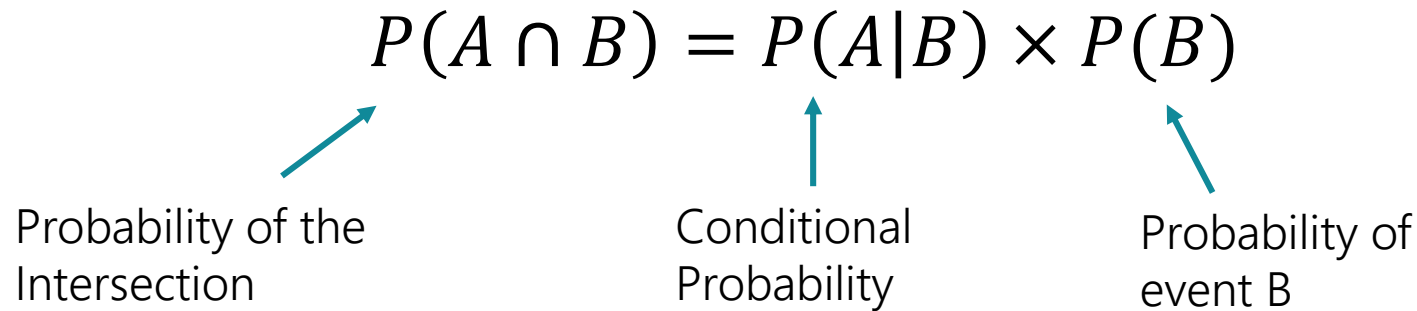
$$P(A \cap B) = P(A|B) \times P(B)$$


Diagram illustrating the Multiplication Rule formula:

- $P(A \cap B)$ is labeled: Probability of the Intersection
- $P(A|B)$ is labeled: Conditional Probability
- $P(B)$ is labeled: Probability of event B

Intuition behind the formula

- We can multiply both sides of the conditional probability formula $P(A|B) = \frac{P(A \cap B)}{P(B)}$ by $P(B)$ to get $P(A \cap B) = P(A|B) \times P(B)$.
- If event B occurs in 40% of the time ($P(B) = 0.4$) and event A occurs in 50% of the time B occurs ($P(A|B) = 0.5$), then they would simultaneously occur in 20% of the time ($P(A|B) \times P(B) = 0.5 \times 0.4 = 0.2$).

Bayes' Law

Bayes' Law helps us understand the relationship between two events by computing the different conditional probabilities. We also call it Bayes' Rule or Bayes' Theorem.

Conditional probability of B, given A.

↓

Conditional probability of A, given B. → $P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$

Intuition behind the formula

- According to the multiplication rule $P(A \cap B) = P(A|B) \times P(B)$, so $P(B \cap A) = P(B|A) \times P(A)$.
- Since $P(A \cap B) = P(B \cap A)$, we plug in $P(B|A) \times P(A)$ for $P(A \cap B)$ in the conditional probability formula $P(A|B) = \frac{P(A \cap B)}{P(B)}$
- Bayes' Law is often used in medical or business analysis to determine which of two symptoms affects the other one more.

An overview of Distributions

A distribution shows the possible values a random variable can take and how frequently they occur.

Important Notation for Distributions:


Y actual outcome

y one of the possible outcomes

$P(Y = y)$ is equivalent to $p(y)$.

We call a function that assigns a probability to each distinct outcome in the sample space, a **probability function**.

	Population	Sample
Mean	μ	\bar{x}
Variance	σ^2	s^2
Standard Deviation	σ	s



Types of Distributions

Certain distributions share characteristics, so we separate them into **types**. The well-defined types of distributions we often deal with have elegant statistics. We distinguish between two big types of distributions based on the type of the possible values for the variable – discrete and continuous.

Discrete



- Have a finite number of outcomes.
- Use formulas we already talked about.
- Can add up individual values to determine probability of an interval.
- Can be expressed with a table, graph or a piece-wise function.
- Expected Values might be unattainable.
- Graph consists of bars lined up one after the other.

Continuous

- Have infinitely many consecutive possible values.
- Use new formulas for attaining the probability of specific values and intervals.
- Cannot add up the individual values that make up an interval because there are **infinitely many** of them.
- Can be expressed with a graph or a continuous function.
- Graph consists of a smooth curve.

Discrete Distributions

Discrete Distributions have finitely many different possible outcomes. They possess several key characteristics which separate them from continuous ones.



Key characteristics of discrete distribution


- Have a finite number of outcomes.
- Use formulas we already talked about.
- Can add up individual values to determine probability of an interval.
- Can be expressed with a table, graph or a piece-wise function.
- Expected Values might be unattainable.
- Graph consists of bars lined up one after the other.
- $P(Y \leq y) = P(Y < y + 1)$

Examples of Discrete Distributions:

- Discrete Uniform Distribution
- Bernoulli Distribution
- Binomial Distribution
- Poisson Distribution

Uniform Distribution

A distribution where all the outcomes are equally likely is called a **Uniform Distribution**.



Notation:

- $Y \sim U(a, b)$
- * alternatively, if the values are categorical, we simply indicate the number of categories, like so: $Y \sim U(a)$

Key characteristics


- All outcomes are equally likely.
- All the bars on the graph are equally tall.
- The expected value and variance have no predictive power.

Example and uses:

- Outcomes of rolling a single die.
- Often used in shuffling algorithms due to its fairness.

Bernoulli Distribution

A distribution consisting of a single trial and only two possible outcomes – success or failure is called a **Bernoulli Distribution**.



Notation:

- $Y \sim \text{Bern}(p)$

Key characteristics


- One trial.
- Two possible outcomes.
- $E(Y) = p$
- $\text{Var}(Y) = p \times (1 - p)$

Example and uses:

- Guessing a single True/False question.
- Often used in when trying to determine what we expect to get out a single trial of an experiment.

Binomial Distribution

A sequence of identical Bernoulli events is called Binomial and follows a **Binomial Distribution**.



Notation:

- $Y \sim B(n, p)$

Key characteristics


- Measures the frequency of occurrence of one of the possible outcomes over the n trials.
- $P(Y = y) = C(y, n) \times p^y \times (1 - p)^{n-y}$
- $E(Y) = n \times p$
- $Var(Y) = n \times p \times (1 - p)$

Example and uses:

- Determining how many times we expect to get a heads if we flip a coin 10 times.
- Often used when trying to predict how likely an event is to occur over a series of trials.

Poisson Distribution

When we want to know the likelihood of a certain event occurring over a given interval of time or distance we use a **Poisson Distribution**.



Notation:

- $Y \sim Po(\lambda)$

Key characteristics


- Measures the frequency over an interval of time or distance. (Only non-negative values.)
- $P(Y = y) = \frac{\lambda^y}{y!e^{-\lambda}}$
- $E(Y) = \lambda$
- $Var(Y) = \lambda$

Example and uses:

- Used to determine how likely a specific outcome is, knowing how often the event **usually** occurs.
- Often incorporated in marketing analysis to determine whether above average visits are out of the ordinary or not.

Continuous Distributions

If the possible values a random variable can take are a sequence of infinitely many consecutive values, we are dealing with a **continuous distribution**.




Key characteristics

- Have infinitely many consecutive possible values.
- Cannot add up the individual values that make up an interval because there are **infinitely many** of them.
- Can be expressed with a graph or a continuous function. Cannot use a table, be
- Graph consists of a smooth curve.
- To calculate the likelihood of an interval, we need integrals.
- They have important CDFs.
- $P(Y = y) = 0$ for any individual value y .
- $P(Y < y) = P(Y \leq y)$

Normal Distribution

A Normal Distribution represents a distribution that most natural events follow.



Notation:

- $Y \sim N(\mu, \sigma^2)$

Key characteristics

- Its graph is bell-shaped curve, symmetric and has thin tails.
- $E(Y) = \mu$
- $Var(Y) = \sigma^2$
- 68% of all its values should fall in the interval:
 - $(\mu - \sigma, \mu + \sigma)$

Example and uses:

- Often observed in the size of animals in the wilderness.
- Could be standardized to use the Z-table.

Standardizing a Normal Distribution

To standardize any normal distribution we need to transform it so that the mean is 0 and the variance and standard deviation are 1.

Using a transformation to create a new random variable z .

$$z = \frac{y - \mu}{\sigma}$$

Ensures mean is 0.


Ensures standard deviation is 1.

Importance of the Standard Normal Distribution.

- The new variable z , represents how many standard deviations away from the mean, each corresponding value is.
- We can transform any Normal Distribution into a Standard Normal Distribution using the transformation shown above.
- Convenient to use because of a table of known values for its CDF, called the Z-score table, or simply the Z-table.

Students' T Distribution

A Normal Distribution represents a small sample size approximation of a Normal Distribution.



Notation:

- $Y \sim t(k)$

Key characteristics


- A small sample size approximation of a Normal Distribution.
- Its graph is bell-shaped curve, symmetric, but has **fat** tails.
- Accounts for extreme values better than the Normal Distribution.
- If $k > 1$: $E(Y) = \mu$ and $Var(Y) = s^2 \times \frac{k}{k-2}$

Example and uses:

- Often used in analysis when examining a small sample of data that usually follows a Normal Distribution.

Chi-Squared Distribution

A Chi-Squared distribution is often used.



Notation:

- $Y \sim \chi^2(k)$

Key characteristics


- Its graph is asymmetric and skewed to the right.
- $E(Y) = k$
- $Var(Y) = 2k$
- The Chi-Squared distribution is the square of the t-distribution.

Example and uses:

- Often used to test goodness of fit.
- Contains a table of known values for its CDF called the χ^2 -table. The only difference is the table shows what part of the table

Exponential Distribution

The **Exponential Distribution** is usually observed in events which significantly change early on.



Notation:

- $Y \sim \text{Exp}(\lambda)$

Key characteristics


- Both the PDF and the CDF plateau after a certain point.
- $E(Y) = \frac{1}{\lambda}$
- $\text{Var}(Y) = \frac{1}{\lambda^2}$
- We often use the natural logarithm to transform the values of such distributions since we do not have a table of known values like the Normal or Chi-Squared.

Example and uses:

- Often used with dynamically changing variables, like online website traffic or radioactive decay.

Logistic Distribution

The **Continuous Logistic Distribution** is observed when trying to determine how continuous variable inputs can affect the probability of a binary outcome.



Notation:

- $Y \sim \text{Logistic}(\mu, s)$

Key characteristics.

- $E(Y) = \mu$
- $\text{Var}(Y) = \frac{s^2 \times \pi^2}{3}$
- The CDF picks up when we reach values near the mean.
- The smaller the scale parameter, the quicker it reaches values close to 1.

Example and uses:

- Often used in sports to anticipate how a player's or team's performance can determine the outcome of the match.

